

# Gaussian EM Algorithm

Greg Johnson

The following is an implementation of the classic EM-algorithm for missing multivariate Gaussian data.

```
X <- matrix(c(NA, 4.605047, 5.8303953, 7.595643, 1.754275, 1.8826819,
4.047683, -1.791576, NA, -1.672295, -3.434457, 2.1768536,
2.904052, -3.906055, -4.6161726), 5, byrow = TRUE)

EM_normal <- function(data, mu0, E0, tol = 1e-05) {

  muhat <- mu0
  Ehat <- E0
  X <- data
  n <- nrow(X)
  p <- ncol(X)
  mre <- 10000 #modified relative error - something large/arbitrary to enter the while loop
  count <- 0
  while (mre > tol) {
    count <- count + 1
    # Step one: Predict
    T1 <- 0
    T2 <- 0
    for (j in 1:n) {
      if (sum(is.na(X[j, ])) == 0) {
        x <- X[j, ]
        T1 <- T1 + x
        T2 <- T2 + x %*% t(x)
        next
      } else {
        mis <- is.na(X[j, ])
        x1 <- X[j, mis] #missing components
        x2 <- X[j, !mis] #not missing
        # we estimate x1 with the conditional expectation of x1 given
        # x2
        muhat1 <- muhat[mis]
        muhat2 <- muhat[!mis]
        Ehat11 <- Ehat[mis, mis]
        Ehat12 <- Ehat[mis, !mis]
        Ehat21 <- Ehat[!mis, mis]
        Ehat22 <- Ehat[!mis, !mis]
        x1est <- muhat1 + Ehat12 %*% solve(Ehat22) %*%
          (x2 - muhat2)
        x <- X[j, ]
        x[mis] <- x1est
        T1 <- T1 + x #est complete data contribution to the sufficient stat T1
        x1x1Test <- Ehat11 - Ehat12 %*% solve(Ehat22) %*%
          Ehat21 + x1est %*% t(x1est)
        xxT <- X[j, ] %*% t(X[j, ])
        xxT[mis, mis] <- x1x1Test
        xxT[!mis, mis] <- x2 %*% t(x1est)
        xxT[mis, !mis] <- x1est %*% t(x2)
      }
    }
  }
}
```

```

        T2 <- T2 + xxT #est complete data contribution to the sufficient stat T2
      }
    }
    # save initial estimate to compute mre
    muhat_old <- muhat
    Ehat_old <- Ehat

    # Step two: Estimate
    muhat <- 1/n * T1
    Ehat <- 1/n * T2 - muhat %*% t(muhat)

    # calculate mre
    par_old <- c(muhat_old, Ehat_old[lower.tri(Ehat_old,
      diag = TRUE)])
    par_new <- c(muhat, Ehat[lower.tri(Ehat, diag = TRUE)])
    mre <- sqrt(sum((par_old - par_new)^2))/max(1, sqrt(sum(par_new^2)))
  }
  return(list(muhat, Ehat))
}

```

```

mu_start <- apply(X, 2, mean, na.rm = TRUE)
Xtemp <- X
Xtemp[1, 1] <- mu_start[1]
Xtemp[3, 3] <- mu_start[3]
E_start <- cov(Xtemp)

EM_normal(X, mu0 = mu_start, E0 = E_start)

```

```

## [[1]]
## [1] 4.4594571 -0.5545532 0.7703368
##
## [[2]]
##      [,1]      [,2]      [,3]
## [1,] 14.930346 11.245574 5.851375
## [2,] 11.245574 10.601760 9.078084
## [3,] 5.851375 9.078084 12.528188

```

Use a modified EM algorithm in which we update  $\Sigma$  from the entire dataset.

```

EM_normal_alt<-function(data,mu0,E0,tol=1e-5){

  muhat<-mu0
  Ehat<-E0
  X<-data
  n<-nrow(X)
  p<-ncol(X)
  mre<-10000 #modified relative error - something large/arbitrary to enter the while loop
  count<-0
  while(mre>tol){
    count<-count+1
    #Step one: Predict
    T1<-0
    X_temp<-X #data with missing data imputed with muhat of current iteration
    #- used to estimate E all at once

```

```

for(j in 1:n){
  if(sum(is.na(X[j,]))==0){
    x<-X[j,]
    T1<-T1+x
    next
  }else{
    mis<-is.na(X[j,])
    x1<-X[j,mis] #missing components
    x2<-X[j,!mis] #not missing
    #we estimate x1 with the conditional expectation of x1 given x2
    muhat1<-muhat[mis]
    muhat2<-muhat[!mis]
    Ehat11<-Ehat[mis,mis]
    Ehat12<-Ehat[mis,!mis]
    Ehat21<-Ehat[!mis,mis]
    Ehat22<-Ehat[!mis,!mis]
    xlest<-muhat1 + Ehat12%*%solve(Ehat22)%*%(x2-muhat2)
    x<-X[j,]
    x[mis]<-xlest
    T1<-T1+x #estimated complete data contribution to the sufficient statistic T1

    X_temp[j,]<-x
  }
}
#save initial estimate to compute mre
muhat_old<-muhat
Ehat_old<-Ehat

#Step two: Estimate
muhat<- 1/n*T1
Ehat<-cov(X_temp)

#calculate mre
par_old<-c(muhat_old,Ehat_old[lower.tri(Ehat_old,diag=TRUE)])
par_new<-c(muhat,Ehat[lower.tri(Ehat,diag=TRUE)])
mre<-sqrt(sum((par_old-par_new)^2))/max(1,sqrt(sum(par_new^2)))
}
return(list(muhat,Ehat))
}

EM_normal_alt(X,mu0=mu_start,E0=E_start)

## [[1]]
## [1] 4.4594644 -0.5545532 0.7703378
##
## [[2]]
##      [,1]      [,2]      [,3]
## [1,] 18.663024 14.05701 7.314264
## [2,] 14.057015 13.25220 11.347603
## [3,] 7.314264 11.34760 15.660230

```

Both algorithms are off the mark in terms of both parameters but they are close considering the small sample size. Both agree on their mle of  $\mu$  but in terms of the mle of  $\Sigma$ , the second method produces an mle estimate that is element-wise larger than the first. With the exception of the variance of  $X_2$ , the elements of sigma are

inflated for both estimates, the second more so. This bias may be a product of the fact that the  $\Sigma$  is not estimated separately and is instead tied to the estimation of  $\mu$ . Based on this simple simulation, I would think that the first method is superior to the second.